

# Unleash the Power within pgBackRest

PostgreSQL User Group NL

Stefan FERCOT

11 March 2021



# Who Am I?

- Stefan Fercot
- aka. pgstef
- <https://pgstef.github.io>
- PostgreSQL user since 2010
- pgBackRest fan & contributor
- Database Backup Architect @EDB

# Agenda

- pgBackRest
  - basic functionalities reminder
  - typical setups
  - less common operations
    - interact with a standby server
    - asynchronous archiving
  - diagnostic
  - monitoring
  - release radar

# pgBackRest

- aims to be a simple, reliable backup and restore system
- written in C (migration completed in 2019)
- custom protocol
  - local or remote operation (via SSH)
- parallelism and asynchronous operation
- full/differential/incremental backup
- backup rotation and archive expiration
- S3/Azure support
- encryption
- multiple compression types (gzip, bzip, lz4, zstd)
- ...

# Installation

- *Use the PGDG repository, Luke!*
  - `yum / apt-get install pgbackrest`

# Configuration

- `/etc/pgbackrest.conf`, example:

```
[global]
repo1-path=/var/lib/pgsql/13/backups
repo1-retention-full=1
log-level-console=info

[my_stanza]
pg1-path=/var/lib/pgsql/13/data
```

- main configuration in the `[global]` part
- each PostgreSQL cluster to backup has its own configuration, called `stanza`

# Setup - archiving

```
# postgresql.conf  
archive_mode = on  
archive_command = 'pgbackrest --stanza=my_stanza archive-push %p'
```

# Initialization

```
$ pgbackrest --stanza=my_stanza stanza-create
P00    INFO: stanza-create command begin 2.32: ...
P00    INFO: stanza-create for stanza 'my_stanza' on repo1
P00    INFO: stanza-create command end: completed successfully

$ pgbackrest --stanza=my_stanza check
P00    INFO: check command begin 2.32: ...
P00    INFO: check repo1 configuration (primary)
P00    INFO: check repo1 archive for WAL (primary)
P00    INFO: WAL segment 0000000100000000000000001 successfully archived to '...' on repo1
P00    INFO: check command end: completed successfully
```



# Full backup

```
$ pgbackrest --stanza=my_stanza --type=full backup
P00    INFO: backup command begin 2.32: ...
P00    INFO: execute non-exclusive pg_start_backup():
backup begins after the next regular checkpoint completes
P00    INFO: backup start archive = 0000000100000000000000003, lsn = 0/3000028
P00    INFO: full backup size = 23.1MB
P00    INFO: execute non-exclusive pg_stop_backup() and wait for all WAL segments to archive
P00    INFO: backup stop archive = 0000000100000000000000003, lsn = 0/3000138
P00    INFO: check archive for segment(s) 0000000100000000000000003:0000000100000000000000003
P00    INFO: new backup label = 20210304-090917F
P00    INFO: backup command end: completed successfully

P00    INFO: expire command begin 2.32: ...
P00    INFO: expire command end: completed successfully
```

# Backup types

- full
  - all database cluster files will be copied
  - no dependencies on previous backups
- incr
  - incremental from the last successful backup
- diff
  - like an incremental backup but always based on the last **full** backup

## INFO command

```
$ pgbackrest info --stanza=my_stanza
stanza: my_stanza
  status: ok
  cipher: none

db (current)
  wal archive min/max (13): 000000010000000000000003/000000010000000000000007

  full backup: 20210304-090917F
    timestamp start/stop: 2021-03-04 09:09:17 / 2021-03-04 09:09:29
    wal start/stop: 000000010000000000000003 / 000000010000000000000003
    database size: 23.1MB, database backup size: 23.1MB
    repo1: backup set size: 2.8MB, backup size: 2.8MB

...

```

# Typical setups

**Do not keep your backup storage on the PostgreSQL host!**

- directly attached storage (`repo1-type`)
- dedicated remote host (`repo1-host`)

# Repository storage types

- `repo1-type`
  - azure - Azure Blob Storage Service
  - cifs - Like posix, but disables links and directory fsyncs
  - posix - Posix-compliant file systems
  - s3 - AWS Simple Storage Service
  - **gcs - Google Cloud Storage coming soon!**

## Dedicated remote host

- create a specific user on the backup server
- install pgBackRest
- setup passwordless SSH connection between the hosts

# Dedicated remote host - configuration

- **pgsql-srv**

```
[global]
repo1-host=backup-srv
repo1-host-user=pgbackrest

[my_stanza]
pg1-path=/var/lib/pgsql/13/data
```

- **backup-srv**

```
[global]
repo1-path=/backup_space

[my_stanza]
pg1-host=pgsql-srv
pg1-host-user=postgres
pg1-path=/var/lib/pgsql/13/data
```

# Command execution with remote storage

- **pgsql-srv**
  - `archive_command`
  - `restore`
- **backup-srv**
  - `backup`



# Less common operations

- refresh Streaming Replication standby
- take backups from the standby server
- asynchronously push or get WAL segments
- selective restore

# Refresh Streaming Replication standby

- repository reachable from both nodes
- add extra stanza configuration on the standby

```
recovery-option=primary_conninfo=host=primary user=replication_user
```

- perform a `delta` restore

```
$ pgbackrest --stanza=my_stanza --type=standby --delta restore
```

- check `primary_conninfo` and `restore_command` before restarting the service

# Take backups from the standby server

- `backup-standby` option

```
[global]
...
backup-standby=y

[my_stanza]
pg1-path=/var/lib/pgsql/13/data
pg2-host=primary
pg2-path=/var/lib/pgsql/13/data
recovery-option=primary_conninfo=host=primary user=replication_user
```

- backup started on primary
  - wait replay location on standby
  - files are copied from the standby

# Asynchronously push WAL segments

- `archive-async=y` and `spool-path`: store temporary data into a local directory
- `archive-push` for `archive_command`
  - write WAL segments into the spool path and acknowledgments when successfully stored in the archive
  - `archive-push-queue-max`:
    - maximum size of the PostgreSQL archive queue
    - prevent the WAL space from filling up until PostgreSQL stops completely...
    - ...but generate missing archives!
- very important to monitor archiving to ensure it continues working

## Asynchronously get WAL segments

- `archive-async=y` and `spool-path`
- `archive-get` for `restore_command`
  - prefetch `archive-get-queue-max` amount of WAL segments to speed up recovery

# Selective restore

- `--db-include` restore option
  - databases not specifically included will be restored as sparse, zeroed files
  - `DROP DATABASE` to remove the zeroed databases after recovery

# Diagnostic

- checksums
- `check` command
- `verify` command

# Checksums

- PostgreSQL `initdb --data-checksums`
  - PGSETUP\_INITDB\_OPTIONS
- `pg_checksums`
  - enable, disable or check data checksums **offline**
- pgBackRest `--checksum-page`
  - validate all data page checksums while backing up a cluster
  - automatically enabled when data page checksums are enabled on the cluster



## `check` command

- validates configuration and `archive_command` setting
- calls `pg_create_restore_point('pgBackRest Archive Check')` and `pg_switch_wal()`

# `verify` command

- internal command only, work in progress

```
pgBackRest 2.32 - 'verify' command help
```

```
Verify the contents of the repository.
```

```
Verify will attempt to determine if the backups and archives in the repository  
are valid.
```

- WAL validation and backup files verification

```
INFO: Results:
```

```
archiveId: 13-1, total WAL checked: 4, total valid WAL: 4
```

```
missing: 0, checksum invalid: 0, size invalid: 0, other: 0
```

```
backup: 20210304-100800F, status: valid, total files checked: 939, total valid files: 939
```

```
missing: 0, checksum invalid: 0, size invalid: 0, other: 0
```

# Monitoring

*Schrödinger's Law of Backups*

*The condition/state of any backup is unknown until a restore is attempted.*

- play with `pgbackrest info --output=json` within PostgreSQL...
- ... or use `check_pgbackrest`

# check\_pgbackrest



- whatever the backups location?
  - only based on `pgbackrest info` output and `repo` commands!

# Available services

```
$ check_pgbbackrest --list
```

```
List of available services:
```

```
archives          Check WAL archives.
```

```
check_pgb_version Check the version of this check_pgbbackrest script.
```

```
retention         Check the retention policy.
```

# Retention

- Fails when
  - the number of full backups is less than `--retention-full`
  - the newest backup is older than `--retention-age`
  - the newest full backup is older than `--retention-age-to-full`

**--retention-full**

```
$ check_pgbackrest --stanza=my_stanza \  
  --service=retention --retention-full=1  
  
BACKUPS_RETENTION OK - backups policy checks ok |  
  full=1 diff=0 incr=0  
  latest=full,20210304-100800F latest_age=163s
```

**--output=human**

```
$ check_pgbackrest --stanza=my_stanza \  
  --service=retention --retention-full=1 --output=human
```

```
Service      : BACKUPS_RETENTION  
Returns      : 0 (OK)  
Message      : backups policy checks ok  
Long message : full=1  
Long message : diff=0  
Long message : incr=0  
Long message : latest=full,20210304-100800F  
Long message : latest_age=3m16s
```



# Multiple arguments together

```
$ check_pgbackrest --stanza=my_stanza \  
  --service=retention --retention-full=1 --output=human \  
  --retention-age=24h --retention-age-to-full=7d
```

```
Service      : BACKUPS_RETENTION  
Returns      : 0 (OK)  
Message      : backups policy checks ok  
Long message : full=1  
Long message : diff=0  
Long message : incr=0  
Long message : latest=full,20210304-100800F  
Long message : latest_age=3m30s  
Long message : latest_full=20210304-100800F  
Long message : latest_full_age=3m30s
```

# Archives

- `info` command
  - shows the oldest (min) archive and the most recent one (max)
  - doesn't check if all the archives in between are really on the disk
  - ...
- `verify` command is still experimental

## Archives (2)

```
$ check_pgbackrest --stanza=my_stanza --service=archives  
  
WAL_ARCHIVES OK - 4 WAL archived, latest archived since 3m21s |  
latest_archive_age=201s num_archives=4
```

**--output=human**

```
$ check_pgbackrest --stanza=my_stanza --service=archives --output=human
```

```
Service           : WAL_ARCHIVES
Returns           : 0 (OK)
Message           : 4 WAL archived
Message           : latest archived since 3m40s
Long message      : latest_archive_age=3m40s
Long message      : num_archives=4
...
Long message      : latest_archive=000000030000000000000000E
Long message      : latest_bck_archive_start=000000030000000000000000B
Long message      : latest_bck_type=full
...
```

# Oops (1)

```
$ rm -rf [...] /archive/my_stanza/13-1/0000000300000000/000000030000000000000000C-*
$ pgbackrest info --stanza=my_stanza
stanza: my_stanza
  status: ok
  cipher: none

  db (current)
    wal archive min/max (13): 000000030000000000000000B/000000030000000000000000E
...
```

```
$ pgbackrest verify --stanza=my_stanza
P00  INFO: Results:
  archiveId: 13-1, total WAL checked: 3, total valid WAL: 3
    missing: 0, checksum invalid: 0, size invalid: 0, other: 0
  backup: 20210304-100800F, status: valid, total files checked: 939, total valid files: 939
    missing: 0, checksum invalid: 0, size invalid: 0, other: 0
```

pgBackRest doesn't report any error!

## Oops (2)

```
$ check_pgbackrest --stanza=my_stanza --service=archives --output=human
Service           : WAL_ARCHIVES
Returns           : 2 (CRITICAL)
Message           : wrong sequence, 1 missing file(s) (...)
Long message      : latest_archive_age=5m49s
Long message      : num_archives=3
Long message      : num_missing_archives=1
Long message      : oldest_missing_archive=000000030000000000000000C
Long message      : latest_missing_archive=000000030000000000000000C
...
```

- WARNING if missing archive < `latest_bck_archive_start`
  - CRITICAL otherwise

# Release Radar

<https://pgbackrest.org/release.html>

## 2.21 (January 15, 2020)

- C migration complete
- Minor bug fixes and improvements
  - 2.22 (January 21, 2020)
  - 2.23 (January 27, 2020)



## 2.24 (February 25, 2020)

- Auto-select backup set for time target

## 2.25 (March 26, 2020)

- LZ4 Compression Support
- `--dry-run` option for the expire command

## 2.26 (April 20, 2020)

- Non-blocking TLS
- TCP keep-alive options

## 2.27 (May 26, 2020)

- Time-based retention for full backup
- Expire a specific backup set
- **Zstandard** and **bzip2** compression support

## 2.28 (July 20, 2020)

- Azure support for repository storage
- `expire-auto` option

## 2.29 (August 31, 2020)

- Automatically retrieve temporary S3 credentials on AWS instances
- `archive-mode` option to disable archiving on restore

## 2.30 (October 5, 2020)

- PostgreSQL 13 support

## 2.31 (December 7, 2020)

- `pg-database` option
- Report page checksums errors in info command text output
  - only when `--set` is provided



## 2.32 (February 8, 2021)

- `repo-ls` and `repo-get` commands

# What's next?

- Multi-repository support
- GCS driver

## Multi-repository

- `archive-push` to multiple repositories
  - very useful for fault-tolerance using `archive-async`
- `backup` to specified location
  - allow multiple retention policies
- `archive-get` able to find archives and history files across all defined repositories
  - very handy in case of gaps - missing archives

# Where

- official website: <https://pgbackrest.org>
- code: <https://github.com/pgbackrest/pgbackrest>
- rpm and deb: in the PGDG repositories!

# User-guides

- official: <https://pgbackrest.org/user-guide.html>
- blog: <https://pgstef.github.io/>

# Conclusion

- pgBackRest is a powerful tool
  - with a lot of features and possibilities
- don't forget *Schrödinger's Law of Backups*
  - monitor backups and archiving system

# Questions?

Thank you for your attention!

